

## Resampling Methods

We are going to use the Auto data to illustrate the results of various resampling methods, so let's load it from the ISLR package and explore.


```
library(ISLR)
data(Auto)
```

```
str(Auto[, -9])
```

```
## 'data.frame':  392 obs. of  8 variables:
## $ mpg      : num  18 15 18 16 17 15 14 14 15 ...
## $ cylinders : num   8  8  8  8  8  8  8  8  8 ...
## $ displacement: num  307 350 318 304 302 429 454 440 455 390 ...
## $ horsepower : num  130 165 150 150 140 198 220 215 225 190 ...
## $ weight      : num 3504 3693 3436 3433 3449 ...
## $ acceleration: num  12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
## $ year        : num  70 70 70 70 70 70 70 70 70 ...
## $ origin      : num   1  1  1  1  1  1  1  1  1 ...
```

A plot is always a nice place to start with a new data set.

```
plot(mpg ~ horsepower, data = Auto)
```



As is exploring the available documentation.

```
?Auto
```

## The Leave-One-Out Cross-Validation (LCOOV) method.

First, lets run a glm model on the Auto data set.

```
glm_auto <- glm(mpg ~ horsepower, data = Auto)
```

Next, load the boot package and check out the documentation for the Cross-validation for Generalized Linear Models function, or cv.glm.

```
library(boot)
```

```
?cv.glm
```

Then, apply cv.glm function to the Auto data set, using glm\_auto model, returning the delta parameter.

```
cv.glm(Auto, glm_auto)$delta
```

```
## [1] 24.23151 24.23114
```

We can speed up the results by writing a function to use the formula displayed in section 5.2 (pg. 180) and then pass the glm\_auto model to it.

```
loocv <- function(x){  
  h <- lm.influence(x)$h  
  mean((residuals(x)/(1-h))^2)  
}
```

Is our new function faster? We can use the system.time function to compare both methods.

```
system.time(  
  cv.glm(Auto, glm_auto)$delta  
)
```

```
##      user  system elapsed  
##      1.3      0.0      1.3
```

```
system.time(  
  loocv(glm_auto)  
)
```


```
##      user  system elapsed  
##       0       0       0
```

Next, let's use a for loop to efficiently create 5 new polynomial versions of the previous model, regressing horsepower against mpg and see if the results improve as polynomial order increases.

```
cv.error <- rep(0, 5)
degree <- 1:5

for(d in degree){
  glm.fit <- glm(mpg ~ poly(horsepower, d), data = Auto)
  cv.error[d] <- loocv(glm.fit)
}


plot(degree, cv.error, type = "b")
```



## The 10-fold Cross-Validation

```
cv.error10 <- rep(0, 5)

for(d in degree){
  glm.fit <- glm(mpg ~ poly(horsepower, d), data = Auto)
  cv.error10[d] <- cv.glm(Auto, glm.fit, K=10)$delta[1]
}
plot(degree, cv.error, type = "b")
lines(degree, cv.error10, type = "b", col = "red")
```



## Bootstrap

Minimum risk investment function from Section 5.2:

```
alpha <- function(x, y) {
  var_x <- var(x)
  var_y <- var(y)
  cov_xy <- cov(x, y)
  (var_y - cov_xy)/(var_x + var_y - 2 * cov_xy)
}

alpha(Portfolio$X, Portfolio$Y)
```

```
## [1] 0.5758321
```

So what is the standard error of alpha?

```
alpha.fn <- function(data, index){  
  with(data[index, ], alpha(X, Y))  
}
```

```
alpha.fn(Portfolio, 1:100)
```

```
## [1] 0.5758321
```

```
set.seed(1)  
alpha.fn(Portfolio, sample(1:100, 100, replace = TRUE))
```

```
## [1] 0.5963833
```

```
boot.out <- boot(Portfolio, alpha.fn, R = 1000)  
boot.out$t0
```

```
## [1] 0.5758321
```

```
plot(boot.out)
```

