

ISLR: Bootstrap quiz

Justin M Shea

Download the file 5.R.RData and load it into R using the load function.

```
data_address <- "https://lagunita.stanford.edu/c4x/HumanitiesSciences/StatLearning/asset/5.R.RData"
download.file(data_address, paste0(getwd(), "/R"))
```

5.R.R1

Consider the linear regression model of y on X_1 and X_2 . What is the standard error for β_1 ?

```
load(path.expand("~/R/Statistical-Learning/data/5.R.RData"))

model_51 <- lm(y ~ X1 + X2, data = Xy)


summary(model_51)

##
## Call:
## lm(formula = y ~ X1 + X2, data = Xy)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.44171 -0.25468 -0.01736  0.33081  1.45860
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.26583    0.01988  13.372 < 2e-16 ***
## X1           0.14533    0.02593   5.604 2.71e-08 ***
## X2           0.31337    0.02923  10.722 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5451 on 997 degrees of freedom
## Multiple R-squared:  0.1171, Adjusted R-squared:  0.1154
## F-statistic: 66.14 on 2 and 997 DF,  p-value: < 2.2e-16
```

5.R.R2

Next, plot the data using `matplot(Xy, type="l")`. Which of the following do you think is most likely given what you see?

```
matplot(Xy, type="l")
```



5.R.R3

Now, use the (standard) bootstrap to estimate s.e. ($\hat{\beta}_1$). To within 10%, what do you get?

```
beta_hat_1 <- function(data, index, formula) {  
  
  model <- lm(formula, data = data[index, ])  
  
  summary(model)$coefficients[2, 1]  
}  
  
library(boot)  
boot_model_51 <- boot(data = Xy, statistic = beta_hat_1, R = 15000, formula = y ~  
  X1 + X2, parallel = "snow", ncpus = 4)  
  
boot.ci(boot_model_51, conf = 0.9)  
  
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS  
## Based on 15000 bootstrap replicates  
##  
## CALL :  
## boot.ci(boot.out = boot_model_51, conf = 0.9)  
##  
## Intervals :  
## Level      Normal          Basic  
## 90%   ( 0.0975,  0.1932 )   ( 0.0969,  0.1932 )  
##  
## Level      Percentile      BCa  
## 90%   ( 0.0974,  0.1937 )   ( 0.0984,  0.1947 )  
## Calculations and Intervals on Original Scale  
  
boot_model_51  
  
##  
## ORDINARY NONPARAMETRIC BOOTSTRAP  
##  
##  
## Call:  
## boot(data = Xy, statistic = beta_hat_1, R = 15000, formula = y ~  
##   X1 + X2, parallel = "snow", ncpus = 4)  
##  
##  
## Bootstrap Statistics :  
##      original      bias    std. error  
## t1* 0.1453263 -2.708126e-05    0.029071
```

5.R.R4

Finally, use the block bootstrap to estimate $\text{s.e.}(\hat{\beta}_1)$. Use blocks of 100 contiguous observations, and resample ten whole blocks with replacement then paste them together to construct each bootstrap time series. For example, one of your bootstrap resamples could be:

```
block_boot_model_51 <- tsboot(Xy, beta_hat_1, formula = y ~ X1 + X2, R = 15000,
  sim = "fixed", l = 100, parallel = "snow", ncpus = 4)

boot.ci(block_boot_model_51, conf = 0.9)
```

```
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 15000 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = block_boot_model_51, conf = 0.9)
##
## Intervals :
## Level      Normal          Basic          Percentile
## 90%   (-0.1802, 0.4742 )  (-0.1746, 0.4735 )  (-0.1829, 0.4652 )
## Calculations and Intervals on Original Scale

block_boot_model_51

##
## BLOCK BOOTSTRAP FOR TIME SERIES
##
## Fixed Block Length of 100
##
## Call:
## tsboot(tseries = Xy, statistic = beta_hat_1, R = 15000, l = 100,
##       sim = "fixed", formula = y ~ X1 + X2, parallel = "snow",
##       ncpus = 4)
##
##
## Bootstrap Statistics :
##      original      bias    std. error
## t1* 0.1453263 -0.001693754  0.1989165
```